

Continues Integration (CI)

Machine Learning Operations

Nicki Skafte Detlefsen,

Postdoc

DTU Compute

Why you should care about today's exercises



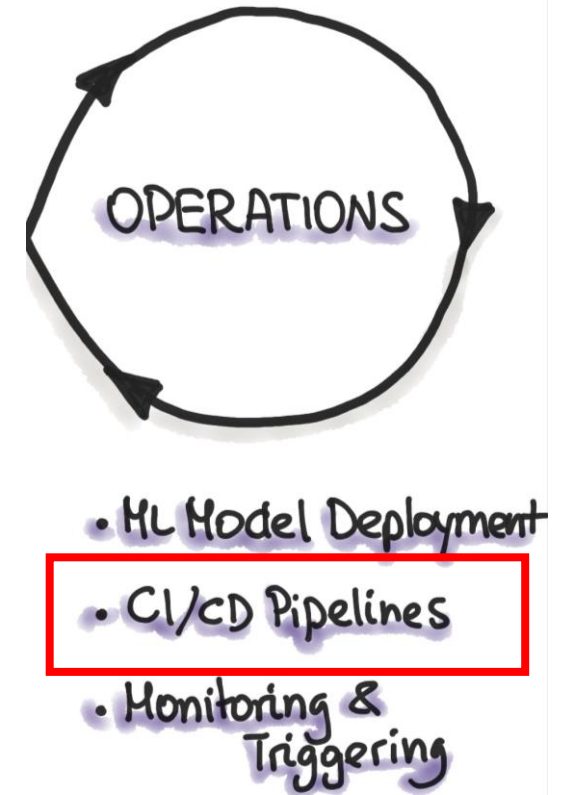
Dev at Fastly : I'll just push this small change to production

Dev at Fastly 2 seconds later:



What is continues integration?

- Software practise
 - Frequently commit code to shared repository
 - By commit sooner than later, errors are captured early
 - Make merging easier
 - Automate build + test
- App independent



What is continues deployment?



- How to get your code to the user
 - 3 phases: Testing, staging and deploying
- App dependent

Note: Not covered in this course



- ML Model Deployment
- CI/CD Pipelines
- Monitoring & Triggering

What should you know about CI?



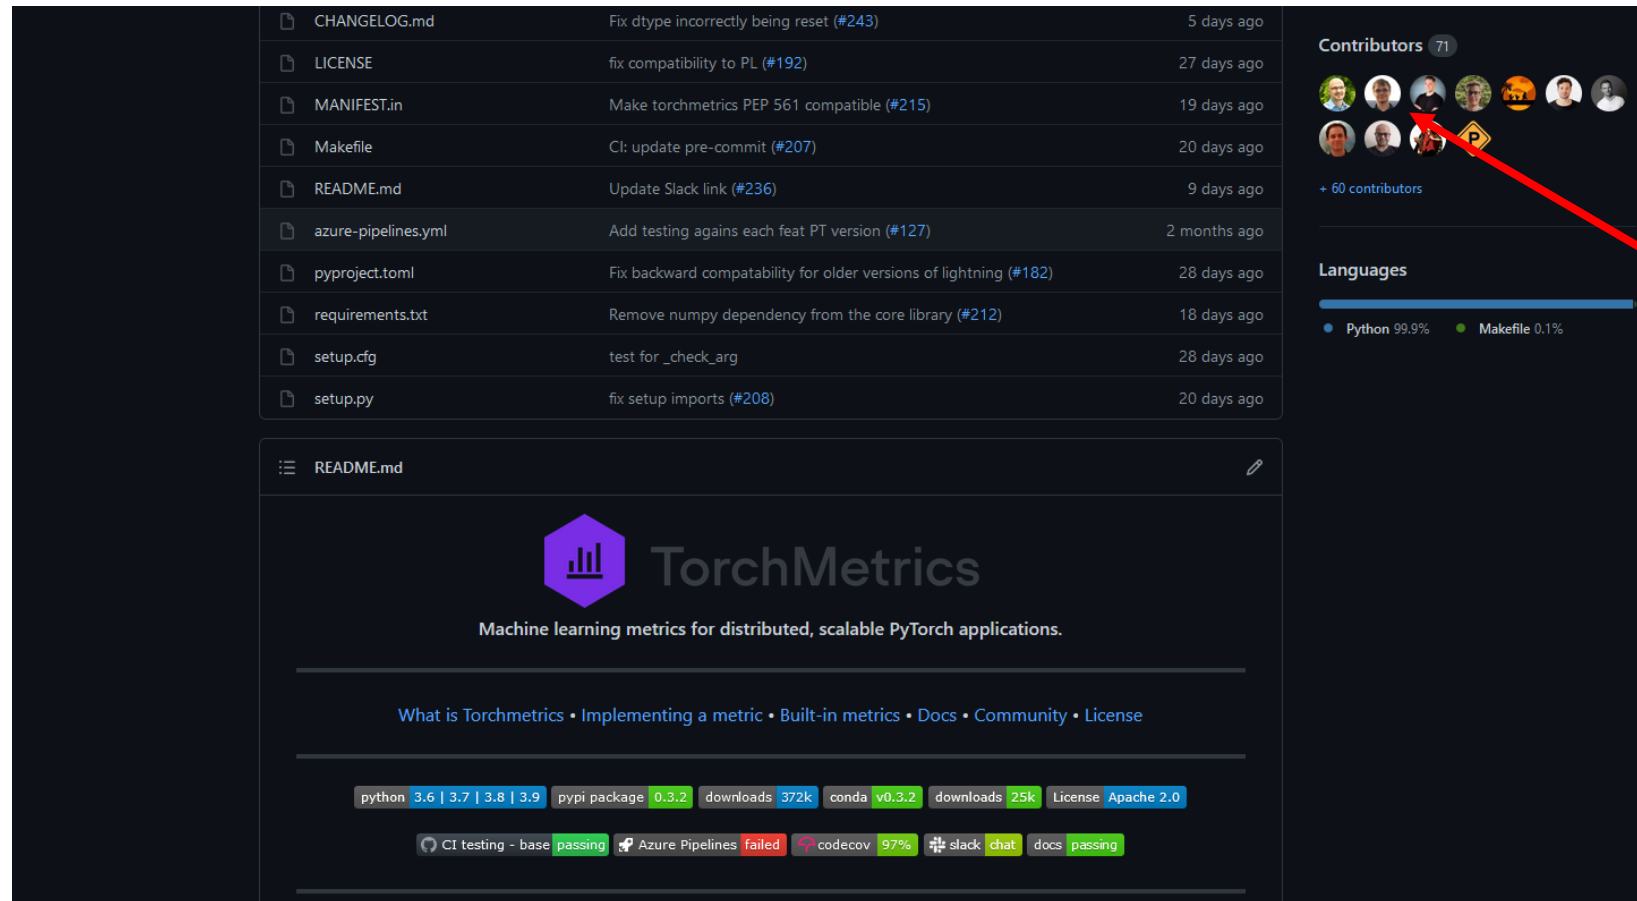
- CI is one of those topics that are best though as learning-by-doing
- If you understand how to use git, the rest is basically googling stuff

Let's look at a practical example

A small case study



All the metric in the world – now in pytorch



File	Description	Time
CHANGELOG.md	Fix dtype incorrectly being reset (#243)	5 days ago
LICENSE	fix compatibility to PL (#192)	27 days ago
MANIFEST.in	Make torchmetrics PEP 561 compatible (#215)	19 days ago
Makefile	CI: update pre-commit (#207)	20 days ago
README.md	Update Slack link (#236)	9 days ago
azure-pipelines.yml	Add testing against each feat PT version (#127)	2 months ago
pyproject.toml	Fix backward compatibility for older versions of lightning (#182)	28 days ago
requirements.txt	Remove numpy dependency from the core library (#212)	18 days ago
setup.cfg	test for _check_arg	28 days ago
setup.py	fix setup imports (#208)	20 days ago

Contributors 71

+ 60 contributors

Languages

- Python 99.9%
- Makefile 0.1%

README.md

TorchMetrics

Machine learning metrics for distributed, scalable PyTorch applications.

What is Torchmetrics • Implementing a metric • Built-in metrics • Docs • Community • License

python 3.6 | 3.7 | 3.8 | 3.9 | pypi package 0.3.2 | downloads 372k | conda v0.3.2 | downloads 25k | License Apache 2.0

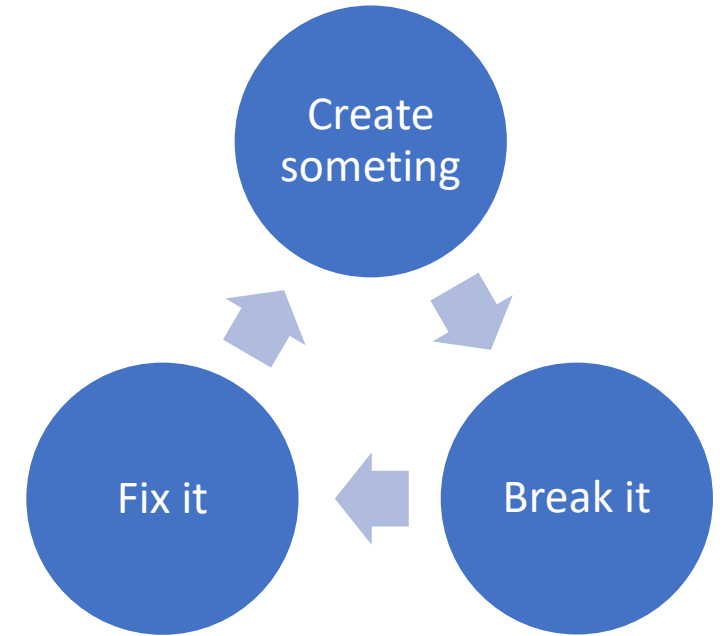
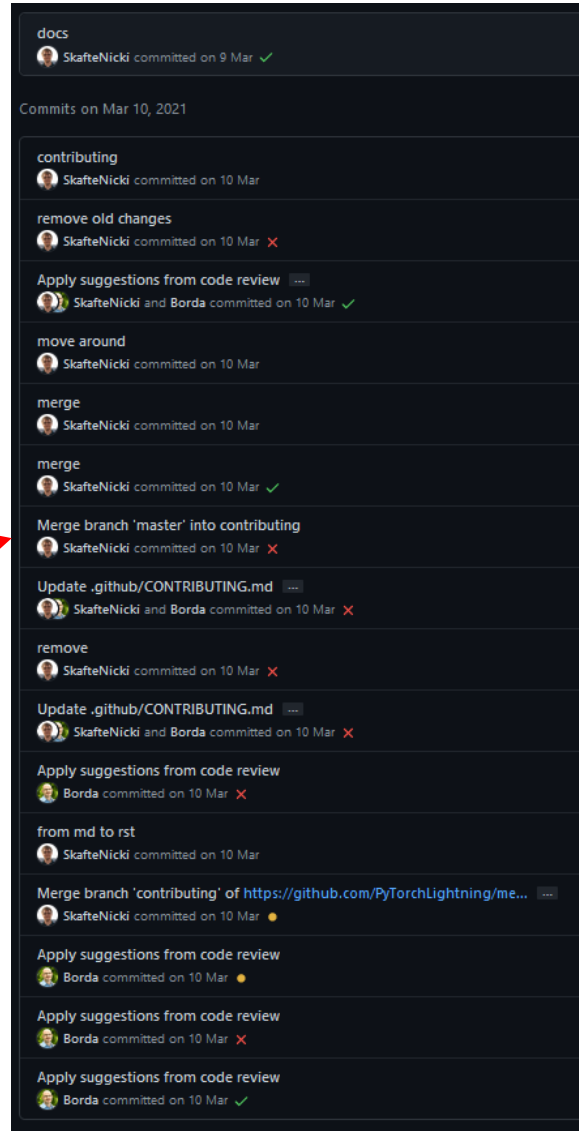
CI testing - base passing | Azure Pipelines failed | codecov 97% | slack chat | docs passing

Have I seen that
guy before?

CI step 1: Committing code



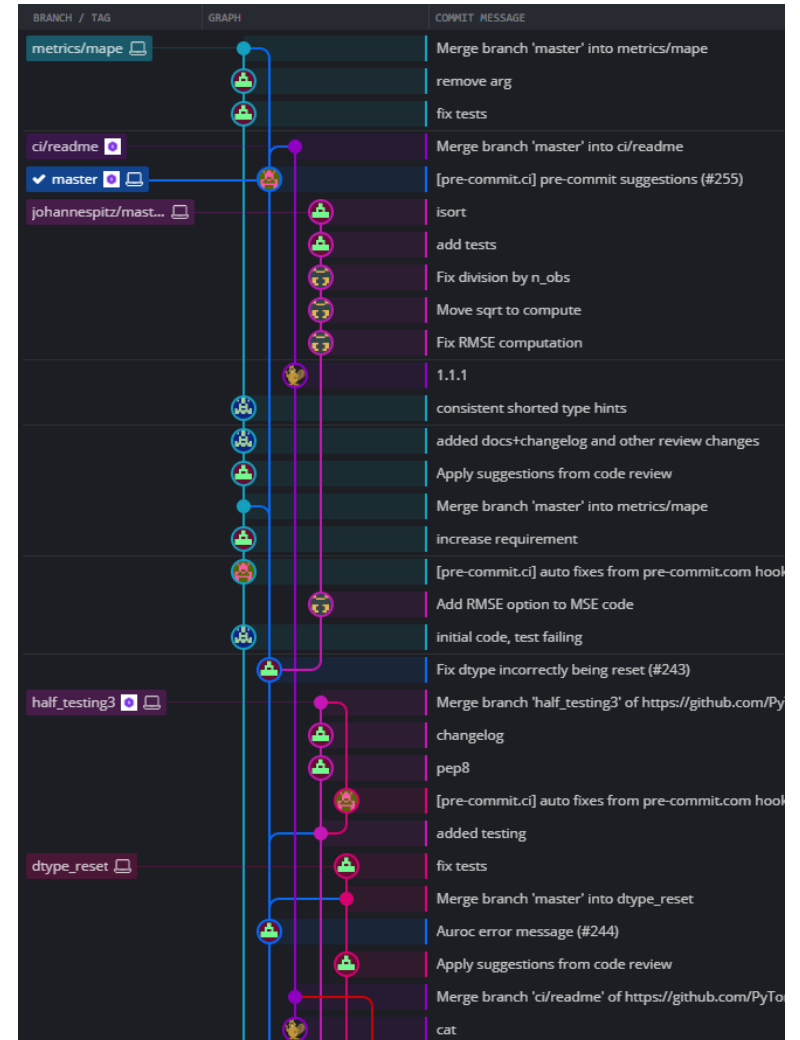
- Commit frequently
 - Catch errors sooner than later
 - Merging can be done automatically



CI step 1: Committing code



- Use branches
 - Enables parallel workflow
 - Experimental features/changes are kept away from "stable" master



CI step 1: Committing code



- Use PRs, other can review your code

A screenshot of a GitHub Pull Request (PR) interface for the repository "PyTorchLightning / metrics". The PR is titled "testing readme examples #140" and is in the "Open" state. The interface shows the PR details, a list of files changed (4 files), and a code diff. Annotations are present: 1. "Find PR" points to the "Pull requests" tab. 2. "Check changed files" points to the "Files changed" tab. 3. "Make one or more comments" points to a comment box where a user has suggested a change to the code. 4. "Send review" points to the "Finish your review" button. The code diff shows changes to the file ".github/set-minimal-versions.py".

1. Find PR

2. Check changed files

3. Make one or more comments

4. Send review

CI step 2: Automating stuff



- What can be automated: EVERYTHING
 - Functional tests
 - Documentation creation
 - Linters (which check style formatting)
 - Security checks
 - Code coverage
 - Custom checks...

A small case study



Ranked by importance (biased)

1. Source code


2. Tests

3. Setup

4. CI workflows


5. Documentation

6. ★



Source code



master	metrics / torchmetrics /	Go to file	Add file	...
	SkaftNicki Fix dtype incorrectly being reset (#243) ...	✓ cda5dbd	5 days ago	History
..				
classification	Allow logit input in classification metrics (#200)		6 days ago	
functional	Fix dtype incorrectly being reset (#243)		5 days ago	
regression	Add differentiability testing to more metrics [3/n] (#225)		13 days ago	
retrieval	Information Retrieval (5/5) (#160)		last month	
utilities	Allow logit input in classification metrics (#200)		6 days ago	
wrappers	Feature pre commit yaml (#145)		2 months ago	
__about__.py	Show must go on (#198)		25 days ago	
__init__.py	Add Specificity metric (#210)		15 days ago	
average.py	allow MetricCollection with args (#176)		29 days ago	
collections.py	Added add_metrics method to MetricCollection (#221)		14 days ago	
metric.py	Fix dtype incorrectly being reset (#243)		5 days ago	
py.typed	Make torchmetrics PEP 561 compatible (#215)		19 days ago	
setup_tools.py	Remove numpy dependency from the core library (#212)		18 days ago	

Test code



In total we have 6365 tests

master	metrics / tests /	Go to file	Add file	...
	SkaftNicki Fix dtype incorrectly being reset (#243) ...	✓ cda5dbd	5 days ago	History
..				
bases	Fix dtype incorrectly being reset (#243)		5 days ago	
classification	Auroc error message (#244)		5 days ago	
functional	Refactor Information Retrieval tests (#156)		last month	
helpers	Ports are reused breaking parallel tests (#226)		14 days ago	
regression	Add differentiability testing to more metrics [3/n] (#225)		13 days ago	
retrieval	Remove numpy dependency from the core library (#212)		18 days ago	
wrappers	Adds indexing operation to Metric class (#142)		2 months ago	
__init__.py	minor refactor tests (#21)		3 months ago	
test_utilities.py	more tests		2 months ago	

Test example 1



Can be simple

```
def test_warning_on_nan(tmpdir):  
    preds = torch.randint(3, size=(20, ))  
    target = torch.randint(3, size=(20, ))  
  
    with pytest.warns(  
        UserWarning,  
        match='.* nan values found in confusion matrix have been replaced with zeros.',  
    ):  
        confusion_matrix(preds, target, num_classes=5, normalize='true')
```

Test example 2



Can be very complex

```
@pytest.mark.parametrize("normalize", ['true', 'pred', 'all', None])
@pytest.mark.parametrize(
    "preds, target, sk_metric, num_classes, multilabel",
    [
        (_input_binary_prob.preds, _input_binary_prob.target, _sk_cm_binary_prob, 2, False),
        (_input_binary_logits.preds, _input_binary_logits.target, _sk_cm_binary_prob, 2, False),
        (_input_binary.preds, _input_binary.target, _sk_cm_binary, 2, False),
        (_input_mlb_prob.preds, _input_mlb_prob.target, _sk_cm_multilabel_prob, NUM_CLASSES, True),
        (_input_mlb_logits.preds, _input_mlb_logits.target, _sk_cm_multilabel_prob, NUM_CLASSES, True),
        (_input_mlb.preds, _input_mlb.target, _sk_cm_multilabel, NUM_CLASSES, True),
        (_input_mcls_prob.preds, _input_mcls_prob.target, _sk_cm_multiclass_prob, NUM_CLASSES, False),
        (_input_mcls_logits.preds, _input_mcls_logits.target, _sk_cm_multiclass_prob, NUM_CLASSES, False),
        (_input_mcls.preds, _input_mcls.target, _sk_cm_multiclass, NUM_CLASSES, False),
        (_input_mdmc_prob.preds, _input_mdmc_prob.target, _sk_cm_multidim_multiclass_prob, NUM_CLASSES, False),
        (_input_mdmc.preds, _input_mdmc.target, _sk_cm_multidim_multiclass, NUM_CLASSES, False)]
)
class TestConfusionMatrix(MetricTester):

    @pytest.mark.parametrize("ddp", [True, False])
    @pytest.mark.parametrize("dist_sync_on_step", [True, False])
    def test_confusion_matrix(
        self, normalize, preds, target, sk_metric, num_classes, multilabel, ddp, dist_sync_on_step
    ):
        self.run_class_metric_test(
            ddp=ddp,
            preds=preds,
            target=target,
            metric_class=ConfusionMatrix,
            sk_metric=partial(sk_metric, normalize=normalize),
            dist_sync_on_step=dist_sync_on_step,
            metric_args={
                "num_classes": num_classes,
                "threshold": THRESHOLD,
                "normalize": normalize,
                "multilabel": multilabel
            }
        )
```

Setup files



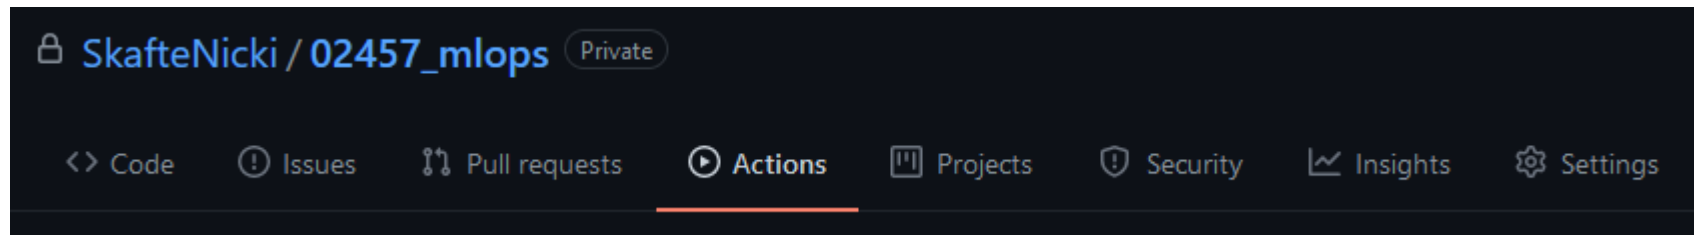
- Contains all information regarding the project
- Allow people to do:
python setup.py install
or if uploaded to pip
pip install my_package

```
setup(
    name='torchmetrics',
    version=about.__version__,
    description=about.__docs__,
    author=about.__author__,
    author_email=about.__author_email__,
    url=about.__homepage__,
    download_url=os.path.join(about.__homepage__, 'archive', 'master.zip'),
    license=about.__license__,
    packages=find_packages(exclude=['tests', 'docs']),
    long_description=long_description,
    long_description_content_type='text/markdown',
    include_package_data=True,
    zip_safe=False,
    keywords=['deep learning', 'machine learning', 'pytorch', 'metrics', 'AI'],
    python_requires='>=3.6',
    setup_requires=[],
    install_requires=setup_tools._load_requirements(_PATH_ROOT),
    project_urls={
        "Bug Tracker": os.path.join(about.__homepage__, 'issues'),
        "Documentation": "https://torchmetrics.rtfd.io/en/latest/",
        "Source Code": about.__homepage__,
    },
    classifiers=[
        'Environment :: Console',
        'Natural Language :: English',
        # How mature is this project? Common values are
        # 3 - Alpha, 4 - Beta, 5 - Production/Stable
        'Development Status :: 3 - Alpha',
        # Indicate who your project is intended for
        'Intended Audience :: Developers',
        'Topic :: Scientific/Engineering :: Artificial Intelligence',
        'Topic :: Scientific/Engineering :: Image Recognition',
        'Topic :: Scientific/Engineering :: Information Analysis',
        # Pick your license as you wish
        # 'License :: OSI Approved :: BSD License',
        'Operating System :: OS Independent',
        # Specify the Python versions you support here. In particular, ensure
        # that you indicate whether you support Python 2, Python 3 or both.
        'Programming Language :: Python :: 3',
        'Programming Language :: Python :: 3.6',
        'Programming Language :: Python :: 3.7',
        'Programming Language :: Python :: 3.8',
        'Programming Language :: Python :: 3.9',
    ],
)
```

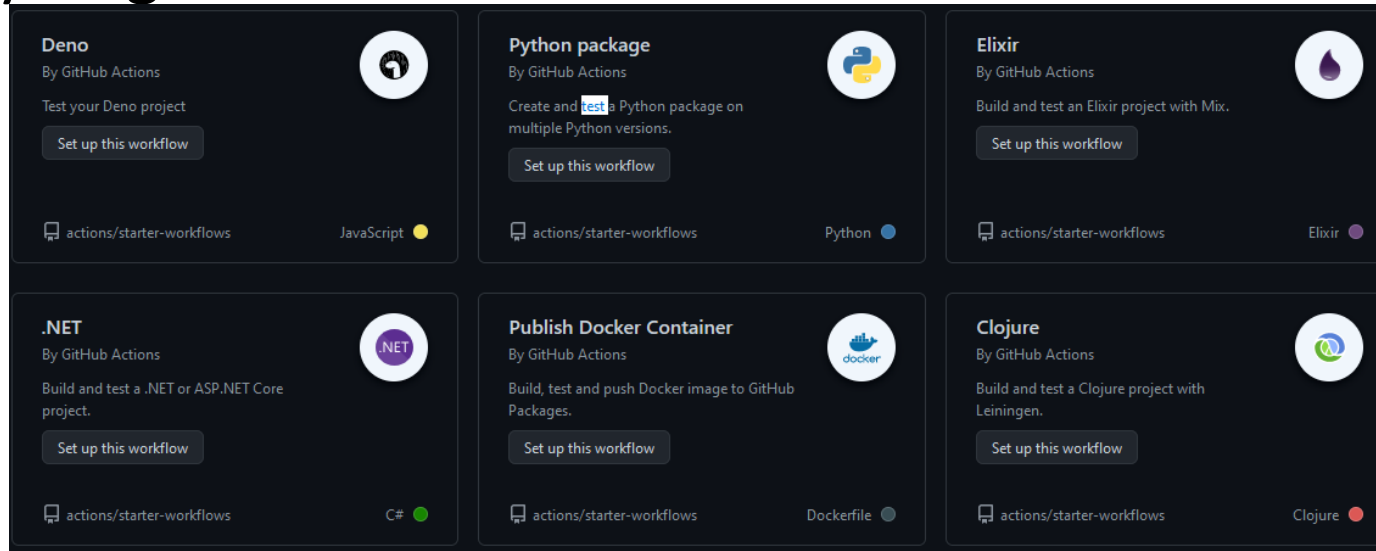

Github actions



- Build-in CI for github
- Free 2,000 automation minutes/month (public repository)



- Many ready to go workflows



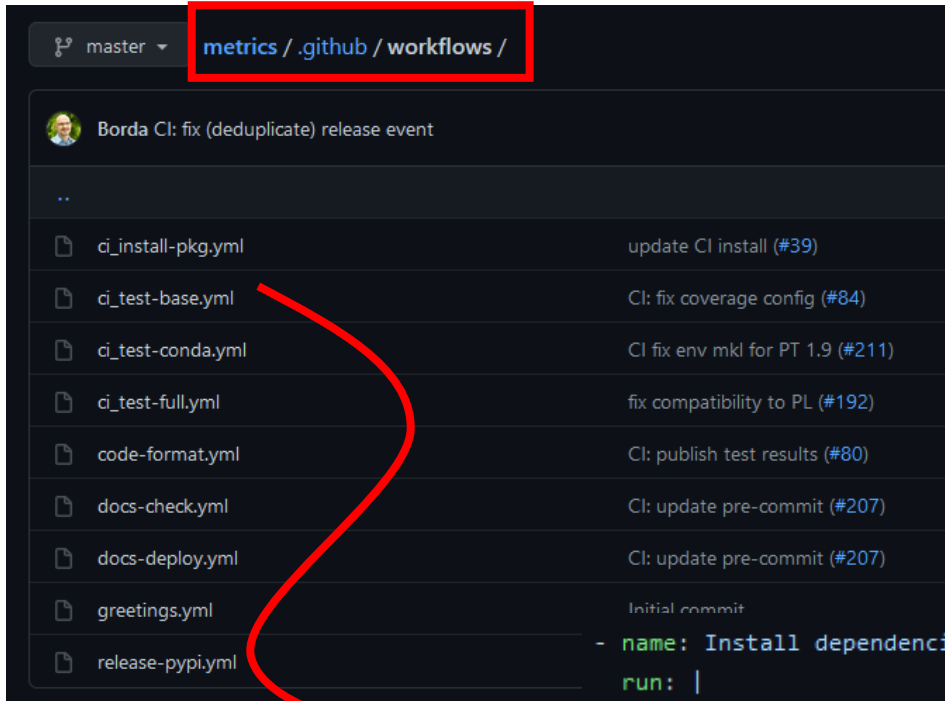
The anatomy of a workflow file



- When tests should be triggered
- Define Operating system + python version
- Setup python
- Install dependencies
- Check linting (stop if errors)
- Run tests

```
02457 mlops / .github / workflows / python-package.yml in main
<> Edit new file Preview
1 name: Python package
2 on:
3   push:
4     branches: [ main ]
5   pull_request:
6     branches: [ main ]
7 jobs:
8   build:
9     runs-on: ubuntu-latest
10    strategy:
11      fail-fast: false
12      matrix:
13        python-version: [3.7, 3.8, 3.9]
14    steps:
15      - uses: actions/checkout@v2
16      - name: Set up Python ${ matrix.python-version }
17        uses: actions/setup-python@v2
18        with:
19          python-version: ${ matrix.python-version }
20      - name: Install dependencies
21        run: |
22          python -m pip install --upgrade pip
23          python -m pip install flake8 pytest
24          if [ -f requirements.txt ]; then pip install -r requirements.txt; fi
25      - name: Lint with flake8
26        run: |
27          # stop the build if there are Python syntax errors or undefined names
28          flake8 . --count --select=E9,F63,F7,F82 --show-source --statistics
29          # exit-zero treats all errors as warnings. The GitHub editor is 127 chars wide
30          flake8 . --count --exit-zero --max-complexity=10 --max-line-length=127 --statistics
31      - name: Test with pytest
32        run: |
33          pytest
```

A small case study



```
Initial commit
- name: Install dependencies
  run: |
    python -m pip install --upgrade --user pip
    pip install --requirement ./requirements.txt --find-links https://download.pytorch.org/whl/cpu/torch_stable.html
    pip install "pytest>6.0" "pytest-cov>2.10" --upgrade-strategy only-if-needed
    python --version
    pip --version
    pip list
  shell: bash

- name: Test Package [only]
  run: |
    # NOTE: run coverage on tests does not propagate faler status for Win, https://github.com/nedbat/coveragepy/issues/1003
    python -m pytest torchmetrics -v --cov=torchmetrics --junitxml=junit/test-results-${{ runner.os }}-${{ matrix.python-version }}.xml
```

All of our workflows – 36 in total



✓ CI testing - complete
on: pull_request

✓ pytest (ubuntu-20.04, 3.6, minimal)

✓ pytest (ubuntu-20.04, 3.8, minimal)

✓ pytest (ubuntu-20.04, 3.8, latest)

✓ pytest (ubuntu-20.04, 3.9, latest)

✓ pytest (macOS-10.15, 3.6, minimal)

✓ pytest (macOS-10.15, 3.8, minimal)

✓ pytest (macOS-10.15, 3.8, latest)

✓ pytest (macOS-10.15, 3.9, latest)

✓ pytest (windows-2019, 3.6, mini...)

✓ pytest (windows-2019, 3.8, mini...)

✓ pytest (windows-2019, 3.8, latest)

✓ pytest (windows-2019, 3.9, latest)

✓ WIP

✓ WIP

✓ Check Code formatting
on: pull_request

✓ flake8

✓ imports-check-isort

✓ typing-check-mypy

✓ Docs check
on: pull_request

✓ test-docs

✓ make-docs

✓ CI testing - base
on: pull_request

✓ doctest (ubuntu-20.04, 3.7)

✓ doctest (windows-2019, 3.7)

✓ doctest (macOS-10.15, 3.7)

✓ Azure Pipelines

✗ PyTorchLightning.metrics Re-run

✓ Mergify

✓ Summary

✓ Codecov

✓ codecov/patch

✓ codecov/project

✓ PyTorch & Conda
on: pull_request

✓ conda (3.7, 1.3)

✓ conda (3.7, 1.4)

✓ conda (3.7, 1.5)

✓ conda (3.7, 1.6)

✓ conda (3.7, 1.7)

⚠ conda (3.7, 1.8)

✓ conda (3.7, 1.9)

Not github actions

✓ Install pkg
on: pull_request

✓ pkg-check

✓ pkg-install (ubuntu-20.04, 3.7)

✓ pkg-install (macOS-10.15, 3.7)

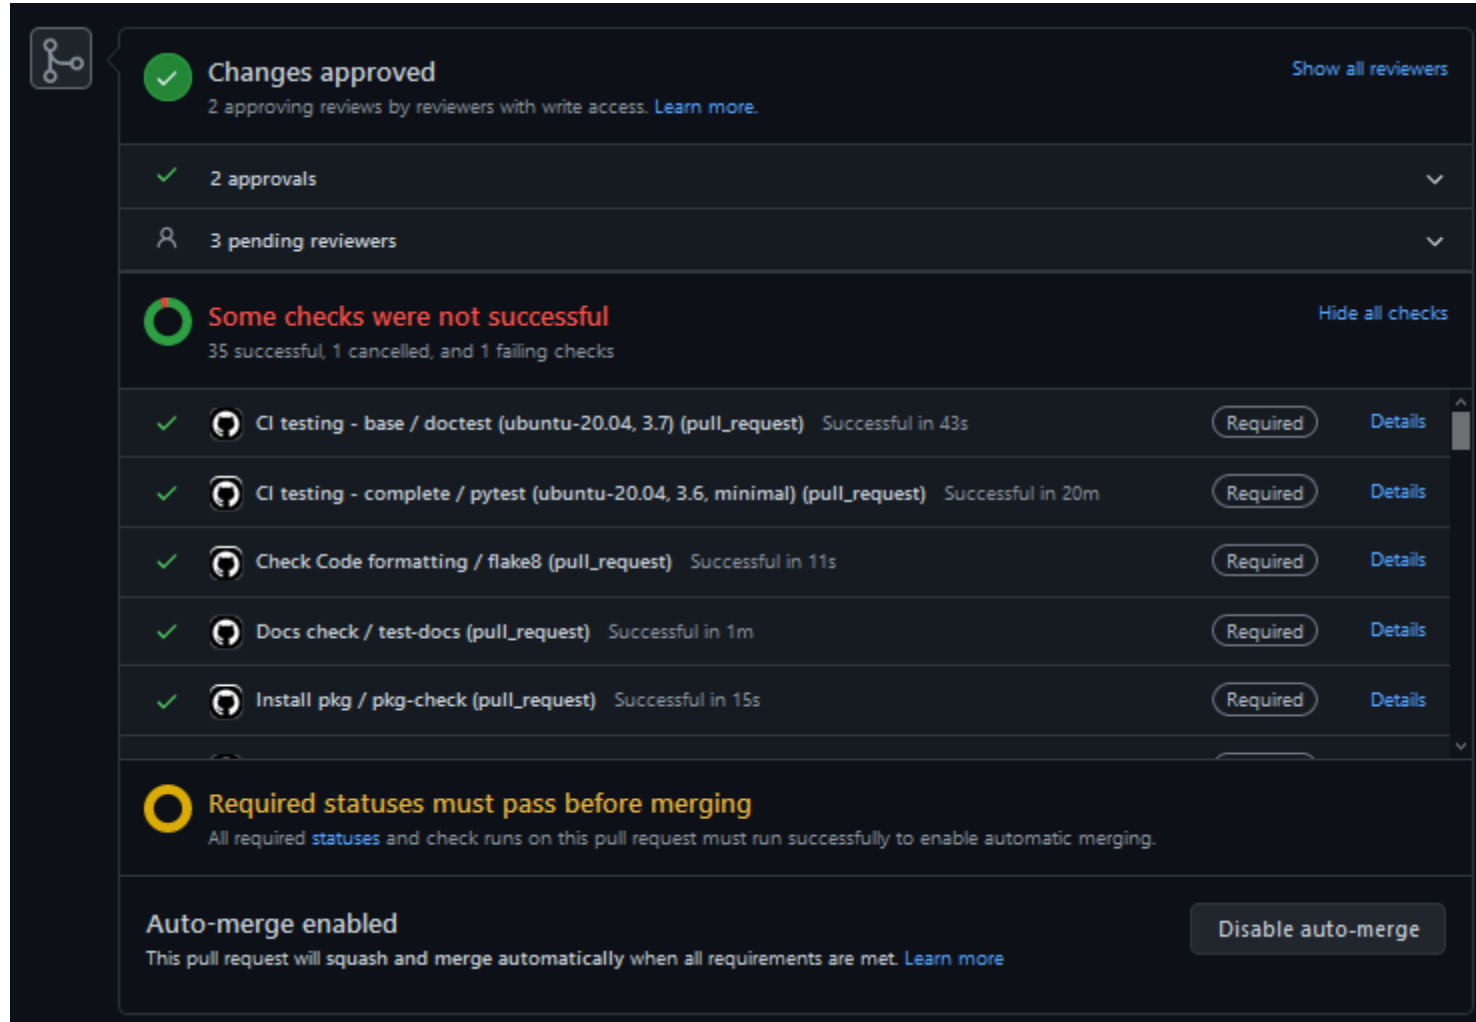
✓ pkg-install (windows-2019, 3.7)

Test combination of

- Hardware setup
- Operating system
- Python version
- Dependencies

Runs tests, docs, coverage, lintint, package install ect

Each PR triggers all tests



The screenshot displays a GitHub pull request interface. At the top, a green checkmark icon indicates 'Changes approved' with the text '2 approving reviews by reviewers with write access. [Learn more.](#)' and a 'Show all reviewers' link. Below this, a summary shows '2 approvals' and '3 pending reviews'. A red circle icon with a slash indicates 'Some checks were not successful' with the text '35 successful, 1 cancelled, and 1 failing checks' and a 'Hide all checks' link. A list of checks follows, each with a green checkmark, a GitHub Actions icon, the check name, duration, and status: 'CI testing - base / doctest (ubuntu-20.04, 3.7) (pull_request)' (Successful in 43s), 'CI testing - complete / pytest (ubuntu-20.04, 3.6, minimal) (pull_request)' (Successful in 20m), 'Check Code formatting / flake8 (pull_request)' (Successful in 11s), 'Docs check / test-docs (pull_request)' (Successful in 1m), and 'Install pkg / pkg-check (pull_request)' (Successful in 15s). Each check has a 'Required' button and a 'Details' link. At the bottom, a yellow circle icon indicates 'Required statuses must pass before merging' with the text 'All required [statuses](#) and check runs on this pull request must run successfully to enable automatic merging.' Below this, it states 'Auto-merge enabled' with the text 'This pull request will squash and merge automatically when all requirements are met. [Learn more](#)' and a 'Disable auto-merge' button.

Changes approved
2 approving reviews by reviewers with write access. [Learn more.](#) [Show all reviewers](#)

2 approvals

3 pending reviews

Some checks were not successful
35 successful, 1 cancelled, and 1 failing checks [Hide all checks](#)

CI testing - base / doctest (ubuntu-20.04, 3.7) (pull_request) Successful in 43s Required Details

CI testing - complete / pytest (ubuntu-20.04, 3.6, minimal) (pull_request) Successful in 20m Required Details

Check Code formatting / flake8 (pull_request) Successful in 11s Required Details

Docs check / test-docs (pull_request) Successful in 1m Required Details

Install pkg / pkg-check (pull_request) Successful in 15s Required Details

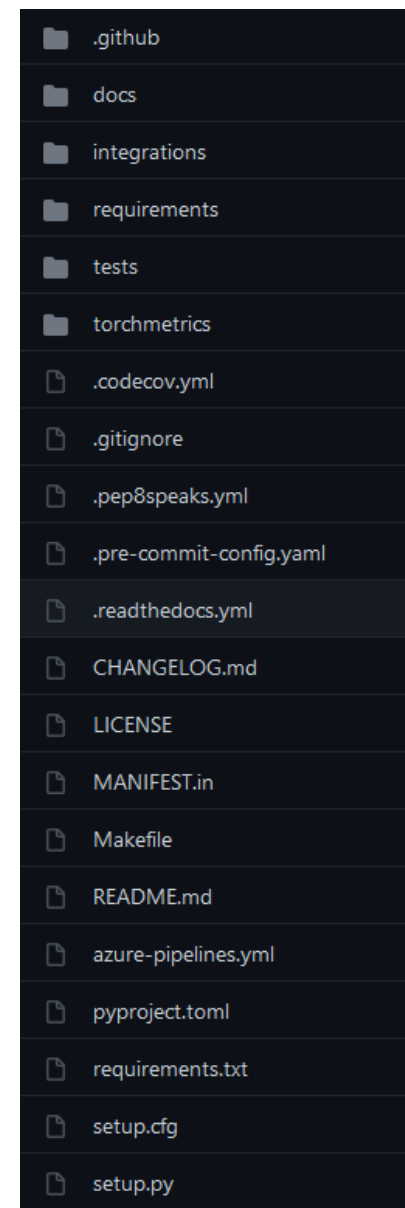
Required statuses must pass before merging
All required [statuses](#) and check runs on this pull request must run successfully to enable automatic merging.

Auto-merge enabled
This pull request will squash and merge automatically when all requirements are met. [Learn more](#) [Disable auto-merge](#)

Other files to consider



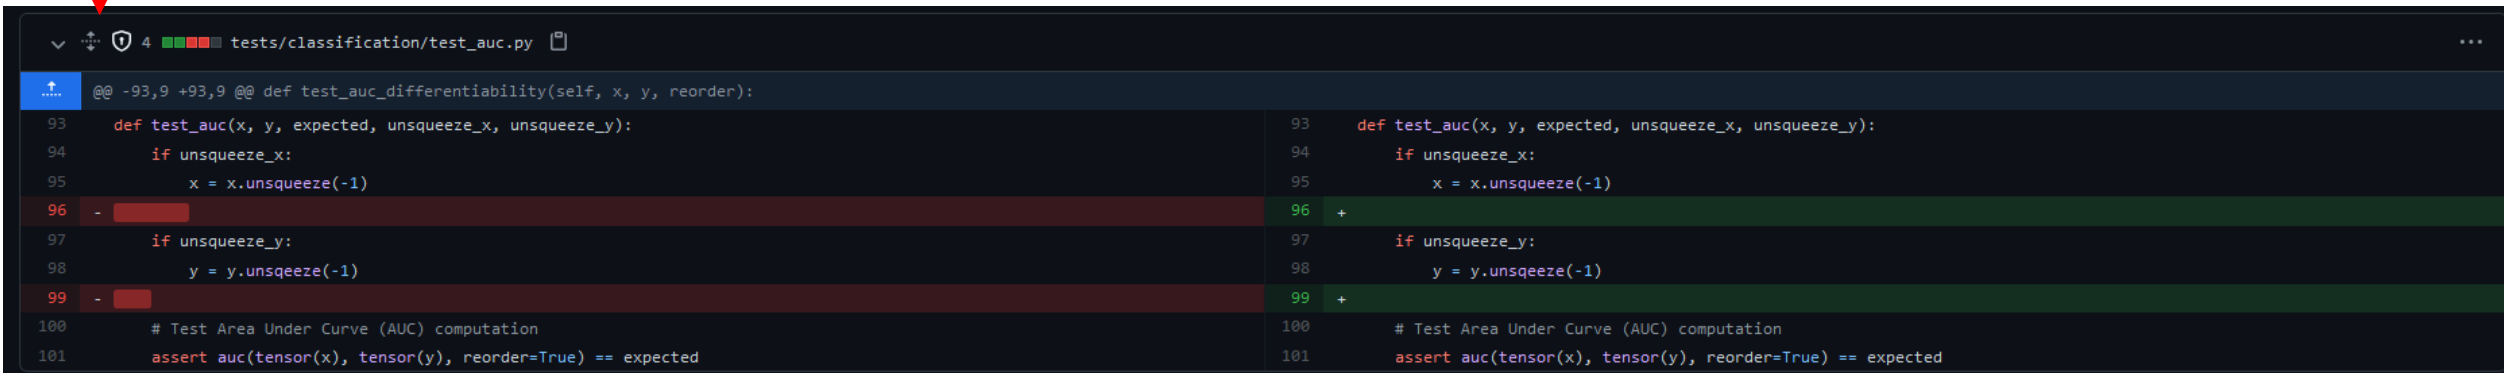
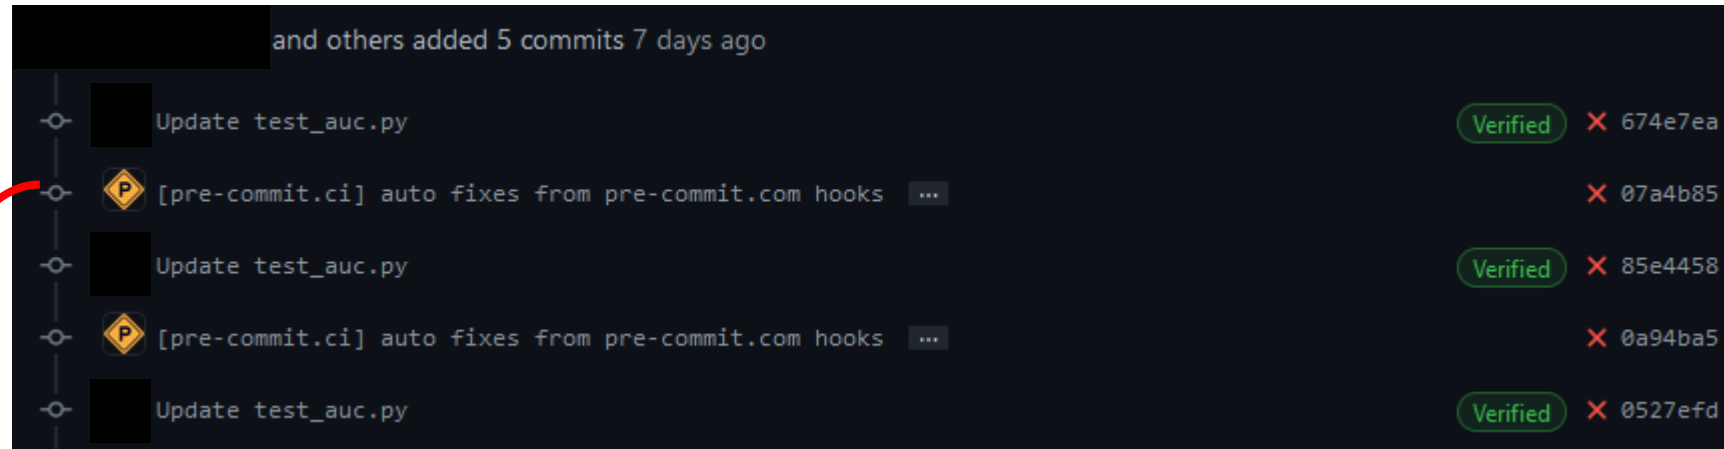
- requirements.txt
- LICENSE
- README.md
- .gitignore



Advanced: Use bots



Bots can take care of tedious task for you (like linting)



The future is here



Gabriele Petronella
@gabro27

So this just happened:

- a bot found a vulnerability in a dependency
- a bot sent a PR to fix it
- the CI verified the PR
- a bot merged it
- a bot celebrated the merge with a GIF

